

# Operation Systems , 2023

## HOMEWORK 3

系級: 資管碩一

學號: 112356045

姓名: 謝東睿

### 1. 說明 clone() 與 fork() 的差異：

fork() 不需要任何參數，僅僅是在創建一個 child process，並為其創建一個獨立於 parent process 的空間。fork() 在創建一個進程時，child process 完全複製 parent process 的資源，複製出來的 child process 有自己的 task\_struct 結構和 pid，但卻複製 parent process 的其它所有資源。

clone() 帶有參數，可以指定創建新的命名空間

(namespace)，還可以將 parent process 的資源有選擇地複製給 child process，而沒有複製的資料結構則通過 pointer 的複製讓 child process 共享，具體要複製哪些資源給 child process，則由參數列表中的 clone\_flags 決定。

### 2. UTS namespace 截圖 x3：

```
[root@m1maclinun:/home/dongrui/os_hw3# hostname  
m1maclinun
```

```
[root@11fe2f3062bf:/# hostname  
11fe2f3062bf
```

```
dongrui@m1maclinunx:~/os_hw3$ hostname  
m1maclinunx
```

### 3. IPC namespace 截圖 x3：

```
root@m1maclinun:/home/dongrui/os_hw3# ipcs
```

```
----- Message Queues -----
```

key	msqid	owner	perms	used-bytes	messages
0xcf7f6888	2	root	644	0	0

```
[root@0359300488ec:/# ipcs
```

```
----- Message Queues -----
```

key	msqid	owner	perms	used-bytes	messages
0x1267e721	0	root	644	0	0

```
root@m1maclinunx:/home/dongrui/os_hw3# ipcs
```

```
----- Message Queues -----
```

key	msqid	owner	perms	used-bytes	messages
0x7b46b9a9	0	root	644	0	0

#### 4. PID namespace 截圖 x3 :

```
root@m1maclinun:/home/dongrui/os_hw3# ls -l /proc/$$/ns
total 0
lrwxrwxrwx 1 root root 0 Jan  5 06:55 cgroup -> 'cgroup:[4026531835]'
lrwxrwxrwx 1 root root 0 Jan  5 08:14 ipc -> 'ipc:[4026531839]'
lrwxrwxrwx 1 root root 0 Jan  5 06:55 mnt -> 'mnt:[4026531841]'
lrwxrwxrwx 1 root root 0 Jan  5 08:14 net -> 'net:[4026531840]'
lrwxrwxrwx 1 root root 0 Jan  5 06:58 pid -> 'pid:[4026531836]'
lrwxrwxrwx 1 root root 0 Jan  5 08:14 pid_for_children -> 'pid:[4026531836]'
lrwxrwxrwx 1 root root 0 Jan  5 08:14 time -> 'time:[4026531834]'
lrwxrwxrwx 1 root root 0 Jan  5 08:14 time_for_children -> 'time:[4026531834]'
lrwxrwxrwx 1 root root 0 Jan  5 08:14 user -> 'user:[4026531837]'
lrwxrwxrwx 1 root root 0 Jan  5 08:14 uts -> 'uts:[4026531838]'
root@0359300488ec:/# ls -l /proc/$$/ns
total 0
lrwxrwxrwx 1 root root 0 Jan  5 08:15 cgroup -> 'cgroup:[4026532488]'
lrwxrwxrwx 1 root root 0 Jan  5 08:15 ipc -> 'ipc:[4026532425]'
lrwxrwxrwx 1 root root 0 Jan  5 08:15 mnt -> 'mnt:[4026532423]'
lrwxrwxrwx 1 root root 0 Jan  5 08:15 net -> 'net:[4026532427]'
lrwxrwxrwx 1 root root 0 Jan  5 08:15 pid -> 'pid:[4026532426]'
lrwxrwxrwx 1 root root 0 Jan  5 08:15 pid_for_children -> 'pid:[4026532426]'
lrwxrwxrwx 1 root root 0 Jan  5 08:15 time -> 'time:[4026531834]'
lrwxrwxrwx 1 root root 0 Jan  5 08:15 time_for_children -> 'time:[4026531834]'
lrwxrwxrwx 1 root root 0 Jan  5 08:15 user -> 'user:[4026531837]'
lrwxrwxrwx 1 root root 0 Jan  5 08:15 uts -> 'uts:[4026532424]'
root@m1maclinunx:/home/dongrui/os_hw3# ls -l /proc/$$/ns
total 0
lrwxrwxrwx 1 root root 0 Jan  5 06:55 cgroup -> 'cgroup:[4026531835]'
lrwxrwxrwx 1 root root 0 Jan  5 08:14 ipc -> 'ipc:[4026531839]'
lrwxrwxrwx 1 root root 0 Jan  5 06:55 mnt -> 'mnt:[4026531841]'
lrwxrwxrwx 1 root root 0 Jan  5 08:14 net -> 'net:[4026531840]'
lrwxrwxrwx 1 root root 0 Jan  5 06:58 pid -> 'pid:[4026531836]'
lrwxrwxrwx 1 root root 0 Jan  5 08:14 pid_for_children -> 'pid:[4026531836]'
lrwxrwxrwx 1 root root 0 Jan  5 08:14 time -> 'time:[4026531834]'
lrwxrwxrwx 1 root root 0 Jan  5 08:14 time_for_children -> 'time:[4026531834]'
lrwxrwxrwx 1 root root 0 Jan  5 08:14 user -> 'user:[4026531837]'
lrwxrwxrwx 1 root root 0 Jan  5 08:14 uts -> 'uts:[4026531838]'
```

#### 5. MNT namespace 截圖 x3 :

root@m1maclinun:/home/dongrui/os\_hw3# ps aux

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.0	0.2	167924	11368	?	Ss	06:55	0:03	/sbin/init
root	2	0.0	0.0	0	0	?	S	06:55	0:00	[kthreadd]
root	3	0.0	0.0	0	0	?	I<	06:55	0:00	[rcu_gp]
root	4	0.0	0.0	0	0	?	I<	06:55	0:00	[rcu_par_gp]
root	5	0.0	0.0	0	0	?	I<	06:55	0:00	[slub_flushwq]
root	6	0.0	0.0	0	0	?	I<	06:55	0:00	[netns]
root	8	0.0	0.0	0	0	?	I<	06:55	0:00	[kworker/0:0H]
root	10	0.0	0.0	0	0	?	I<	06:55	0:00	[mm_percpu_wq]
root	11	0.0	0.0	0	0	?	S	06:55	0:00	[rcu_tasks_ru]
root	12	0.0	0.0	0	0	?	S	06:55	0:00	[rcu_tasks_tr]
root	13	0.0	0.0	0	0	?	S	06:55	0:00	[ksoftirqd/0]
root	14	0.0	0.0	0	0	?	I	06:55	0:00	[rcu_sched]
root	15	0.0	0.0	0	0	?	S	06:55	0:00	[migration/0]
root	16	0.0	0.0	0	0	?	S	06:55	0:00	[idle_inject/
root	17	0.0	0.0	0	0	?	I	06:55	0:02	[kworker/0:1-
root	18	0.0	0.0	0	0	?	S	06:55	0:00	[cpuhp/0]
root	19	0.0	0.0	0	0	?	S	06:55	0:00	[cpuhp/1]
root	20	0.0	0.0	0	0	?	S	06:55	0:00	[idle_inject/
root	21	0.0	0.0	0	0	?	S	06:55	0:00	[migration/1]
root	22	0.0	0.0	0	0	?	S	06:55	0:00	[ksoftirqd/1]
root	24	0.0	0.0	0	0	?	I<	06:55	0:00	[kworker/1:0H]
root	25	0.0	0.0	0	0	?	S	06:55	0:00	[cpuhp/2]
root	26	0.0	0.0	0	0	?	S	06:55	0:00	[idle_inject/
root	27	0.0	0.0	0	0	?	S	06:55	0:00	[migration/2]
root	28	0.0	0.0	0	0	?	S	06:55	0:00	[ksoftirqd/2]
root	30	0.0	0.0	0	0	?	I<	06:55	0:00	[kworker/2:0H]
root	31	0.0	0.0	0	0	?	S	06:55	0:00	[cpuhp/3]
root	32	0.0	0.0	0	0	?	S	06:55	0:00	[idle_inject/
root	33	0.0	0.0	0	0	?	S	06:55	0:00	[migration/3]
root	34	0.0	0.0	0	0	?	S	06:55	0:00	[ksoftirqd/3]
root	35	0.0	0.0	0	0	?	I	06:55	0:03	[kworker/3:0-
root	36	0.0	0.0	0	0	?	I<	06:55	0:00	[kworker/3:0H]
root	37	0.0	0.0	0	0	?	S	06:55	0:00	[kdevtmpfs]
root	38	0.0	0.0	0	0	?	I<	06:55	0:00	[inet_frag_wq]
root	40	0.0	0.0	0	0	?	I	06:55	0:01	[kworker/1:1-
root	42	0.0	0.0	0	0	?	S	06:55	0:00	[kauditd]
root	43	0.0	0.0	0	0	?	S	06:55	0:00	[khungtaskd]
root	44	0.0	0.0	0	0	?	S	06:55	0:00	[oom_reaper]
root	45	0.0	0.0	0	0	?	I<	06:55	0:00	[writeback]
root	46	0.0	0.0	0	0	?	S	06:55	0:00	[kcompactd0]
root	47	0.0	0.0	0	0	?	SN	06:55	0:00	[ksmd]
root	48	0.0	0.0	0	0	?	SN	06:55	0:00	[khugepaged]
root	94	0.0	0.0	0	0	?	I<	06:55	0:00	[kintegrityd]
root	95	0.0	0.0	0	0	?	I<	06:55	0:00	[kblockd]
root	96	0.0	0.0	0	0	?	I<	06:55	0:00	[blkcg_punt_b]
root	97	0.0	0.0	0	0	?	I<	06:55	0:00	[tpm_dev_wq]
root	98	0.0	0.0	0	0	?	I<	06:55	0:00	[ata_sff]
root	99	0.0	0.0	0	0	?	I<	06:55	0:00	[md]
root	100	0.0	0.0	0	0	?	I<	06:55	0:00	[edac-poller]
root	101	0.0	0.0	0	0	?	I<	06:55	0:00	[devfreq_wq]
root	102	0.0	0.0	0	0	?	S	06:55	0:00	[watchdogd]
root	104	0.0	0.0	0	0	?	I<	06:55	0:00	[kworker/1:1H]
root	105	0.0	0.0	0	0	?	S	06:55	0:00	[kswapd0]
root	106	0.0	0.0	0	0	?	S	06:55	0:00	[ecryptfs-kth
root	108	0.0	0.0	0	0	?	I<	06:55	0:00	[kthrotld]
root	109	0.0	0.0	0	0	?	S	06:55	0:00	[irq/49-ACPI:
root	110	0.0	0.0	0	0	?	I<	06:55	0:00	[acpi_thermal
root	112	0.0	0.0	0	0	?	I<	06:55	0:00	[mld]
root	113	0.0	0.0	0	0	?	I<	06:55	0:00	[ipv6_addrcon

root@0359300488ec:/# ps aux

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.0	0.0	4136	3408	pts/0	Ss	08:10	0:00	/bin/bash
root	18	0.0	0.0	6412	1648	pts/0	R+	08:28	0:00	ps aux

```

dongrui 1534 0.0 0.1 106300 5652 ? S 06:55 0:00 (sd-pam)
dongrui 1541 0.0 0.1 8256 5088 tty1 S 06:55 0:00 -bash
root 1608 0.0 0.4 297552 18840 ? Ssl 06:57 0:00 /usr/libexec/packagekitd
root 2294 0.0 1.0 1944972 43580 ? Ssl 06:58 0:14 /usr/bin/containerd
root 2532 0.1 1.8 1591800 72112 ? Ssl 06:58 0:05 /usr/bin/dockerd -H fd:// --containerd=/run/containerd/containerd.sock
root 2877 0.0 0.0 0 0 ? I 07:00 0:00 [kworker/2:0-events]
root 3532 0.0 0.0 0 0 ? I 07:10 0:00 [kworker/3:1]
root 3545 0.0 0.2 18400 9776 ? Ss 07:28 0:00 sshd: dongrui [priv]
dongrui 3600 0.0 0.1 18532 6532 ? S 07:28 0:00 sshd: dongrui@pts/0
dongrui 3601 0.0 0.1 8244 5060 pts/0 Ss 07:28 0:00 -bash
root 3623 0.0 0.0 0 0 ? I 07:42 0:00 [kworker/u8:0-events_power_efficient]
root 3639 0.0 0.1 13020 4788 pts/0 S+ 07:46 0:00 sudo ./uts
root 3640 0.0 0.0 13020 1840 pts/1 Ss 07:46 0:00 sudo ./uts
root 3641 0.0 0.0 3212 764 pts/1 S 07:46 0:00 ./uts
root 3642 0.0 0.0 7144 3836 pts/1 S 07:46 0:00 /bin/bash
root 3812 0.0 0.0 0 0 ? I 07:55 0:01 [kworker/u8:3-events_unbound]
root 3833 0.0 0.1 13400 4780 pts/1 S+ 08:00 0:00 sudo ./ipc
root 3834 0.0 0.0 13400 1836 pts/4 Ss 08:00 0:00 sudo ./ipc
root 3835 0.0 0.0 3212 752 pts/4 S 08:00 0:00 ./ipc
root 3836 0.0 0.0 7144 3840 pts/4 S 08:00 0:00 /bin/bash
root 3878 0.0 0.1 12988 4788 tty1 R+ 08:05 0:00 sudo ./ipc
root 3879 0.0 0.0 12988 676 pts/5 Ss 08:05 0:00 sudo ./ipc
root 3880 0.0 0.0 3212 760 pts/5 S 08:05 0:00 ./ipc
root 3881 0.0 0.0 7144 3860 pts/5 S 08:05 0:00 /bin/bash
root 3903 0.0 0.0 0 0 ? R 08:08 0:00 [kworker/u8:1-events_unbound]
root 3910 0.0 0.0 0 0 ? I 08:09 0:00 [kworker/2:1-kdmflush]
root 3936 0.0 0.0 0 0 ? R 08:09 0:00 [kworker/1:0-events]
root 3952 0.0 0.2 18400 9648 ? Ss 08:09 0:00 sshd: dongrui [priv]
dongrui 4048 0.0 0.1 18532 6524 ? S 08:09 0:00 sshd: dongrui@pts/2
dongrui 4049 0.0 0.1 8244 5068 pts/2 Ss 08:09 0:00 -bash
root 4058 0.0 0.1 13020 4784 pts/2 S+ 08:10 0:00 sudo docker run -it ubuntu /bin/bash
root 4059 0.0 0.0 13020 1836 pts/3 Ss 08:10 0:00 sudo docker run -it ubuntu /bin/bash
root 4060 0.0 0.5 1401236 23288 pts/3 S1+ 08:10 0:00 docker run -it ubuntu /bin/bash
root 4100 0.0 0.3 721968 12720 ? Sl 08:10 0:00 /usr/bin/containerd-shim-runc-v2 -namespace moby -id 0359300480ec0f878a38235785a487d1781cc703
root 4120 0.0 0.0 4136 3408 pts/0 Ss+ 08:10 0:00 /bin/bash
root 4159 0.0 0.1 12984 4776 pts/4 S+ 08:14 0:00 sudo ./pid
root 4160 0.0 0.0 12984 676 pts/5 Ss 08:14 0:00 sudo ./pid
root 4161 0.0 0.0 3212 768 pts/6 S 08:14 0:00 ./pid
root 4162 0.0 0.0 7144 3852 pts/6 S+ 08:14 0:00 /bin/bash
root 4174 0.0 0.0 0 0 ? I 08:16 0:00 [kworker/0:0-events]
root 4175 0.0 0.0 5592 800 tty2 Ss+ 08:16 0:00 /sbin/agetty -o -p -- \u --noclear tty2 linux
root 4184 0.0 0.1 13400 4788 pts/5 S+ 08:17 0:00 sudo ./pid
root 4185 0.0 0.0 13400 1832 pts/7 Ss 08:17 0:00 sudo ./pid
root 4186 0.0 0.0 3212 752 pts/7 S 08:17 0:00 ./pid
root 4187 0.0 0.0 7144 3856 pts/7 S 08:17 0:00 /bin/bash
root 4201 0.1 0.0 0 0 ? I 08:23 0:00 [kworker/u8:2-ext4-fsv-conversion]
root 4210 0.0 0.0 0 0 ? I 08:24 0:00 [kworker/0:2-events]
root 4246 0.0 0.0 0 0 ? I 08:26 0:00 [kworker/2:2]
root 4253 0.0 0.0 9420 1652 pts/7 R+ 08:26 0:00 ps aux

```

## 6. Docker 上使用資源截圖 截圖 x2：

CONTAINER ID	NAME	CPU %	MEM USAGE / LIMIT	MEM %	NET I/O
5020fe3ba23c	stress	90.64%	5.742MiB / 3.817GiB	0.15%	726B /
0B	0B / 0B				2
18d64b321a5a	musng_rhodes	0.00%	820KiB / 3.817GiB	0.02%	1.46kB
/ 0B	0B / 0B				1
█					
CONTAINER ID	NAME	CPU %	MEM USAGE / LIMIT	MEM %	NET I/O
6ed4b9f80fa1	stress	20.88%	5.738MiB / 3.817GiB	0.15%	796B /
0B	0B / 0B				2
18d64b321a5a	musng_rhodes	0.00%	820KiB / 3.817GiB	0.02%	1.53kB
/ 0B	0B / 0B				1
█					

## 7. 使用 Cgroup 操做 deadloop.c 的資源使用 截圖 x2：



```
top - 10:14:15 up 33 min, 4 users, load average: 0.69, 0.71, 0.63
Tasks: 110 total, 2 running, 108 sleeping, 0 stopped, 0 zombie
%Cpu(s): 99.3 us, 0.7 sy, 0.0 ni, 0.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
MiB Mem : 3909.0 total, 2205.0 free, 274.5 used, 1429.5 buff/cache
MiB Swap: 3908.0 total, 3908.0 free, 0.0 used. 3375.7 avail Mem

  PID USER      PR  NI  VIRT  RES  SHR S %CPU  %MEM    TIME+  COMMAND
 3846 root        20   0   2640   956   868 R  98.0   0.0   0:46.37 deadloop
 3887 dongrui    20   0  10472  3692  3136 R   0.7   0.1   0:00.21 top
 3889 root        20   0     0     0     0 I   0.7   0.0   0:00.02 kworker+
 2484 root        20   0 1211232 49048 34744 S   0.3   1.2   0:06.52 contain+
 3231 dongrui   20   0  17300   8016  5608 S   0.3   0.2   0:00.53 sshd
 3261 root        20   0     0     0     0 I   0.3   0.0   0:01.62 kworker+
   1 root        20   0 167676 13004  8156 S   0.0   0.3   0:33.44 systemd
   2 root        20   0     0     0     0 S   0.0   0.0   0:00.03 kthreadd
   3 root         0 -20     0     0     0 I   0.0   0.0   0:00.00 rcu_gp
   4 root         0 -20     0     0     0 I   0.0   0.0   0:00.00 rcu_par+
   5 root         0 -20     0     0     0 I   0.0   0.0   0:00.00 slub_fl+
   6 root         0 -20     0     0     0 I   0.0   0.0   0:00.00 netns
   7 root        20   0     0     0     0 I   0.0   0.0   0:06.11 kworker+
   8 root         0 -20     0     0     0 I   0.0   0.0   0:00.00 kworker+
  10 root         0 -20     0     0     0 I   0.0   0.0   0:00.00 mm_perc+
  11 root        20   0     0     0     0 S   0.0   0.0   0:00.00 rcu_tas+
  12 root        20   0     0     0     0 S   0.0   0.0   0:00.00 rcu_tas+

top - 10:21:05 up 40 min, 6 users, load average: 0.60, 0.81, 0.73
Tasks: 113 total, 2 running, 111 sleeping, 0 stopped, 0 zombie
%Cpu(s): 20.7 us, 1.3 sy, 0.0 ni, 78.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
MiB Mem : 3909.0 total, 2203.8 free, 275.7 used, 1429.6 buff/cache
MiB Swap: 3908.0 total, 3908.0 free, 0.0 used. 3374.6 avail Mem

  PID USER      PR  NI  VIRT  RES  SHR S %CPU  %MEM    TIME+  COMMAND
 3846 root        20   0   2640   956   868 R  19.7   0.0   6:15.26 deadloop
 3982 dongrui    20   0  10472  3728  3172 R   1.6   0.1   0:00.58 top
 3261 root        20   0     0     0     0 I   0.7   0.0   0:03.34 kworker+
  25 root        20   0     0     0     0 S   0.3   0.0   0:00.91 kcompac+
 2484 root        20   0 1211232 48932 34808 S   0.3   1.2   0:07.61 contain+
 3231 dongrui   20   0  17300   8016  5608 S   0.3   0.2   0:01.27 sshd
   1 root        20   0 167676 13004  8156 S   0.0   0.3   0:33.53 systemd
   2 root        20   0     0     0     0 S   0.0   0.0   0:00.03 kthreadd
   3 root         0 -20     0     0     0 I   0.0   0.0   0:00.00 rcu_gp
   4 root         0 -20     0     0     0 I   0.0   0.0   0:00.00 rcu_par+
   5 root         0 -20     0     0     0 I   0.0   0.0   0:00.00 slub_fl+
   6 root         0 -20     0     0     0 I   0.0   0.0   0:00.00 netns
   8 root         0 -20     0     0     0 I   0.0   0.0   0:00.00 kworker+
  10 root         0 -20     0     0     0 I   0.0   0.0   0:00.00 mm_perc+
  11 root        20   0     0     0     0 S   0.0   0.0   0:00.00 rcu_tas+
  12 root        20   0     0     0     0 S   0.0   0.0   0:00.00 rcu_tas+
  13 root        20   0     0     0     0 S   0.0   0.0   0:02.31 ksofttir+
```

## 8. Runc - 使用 ps aux 指令列出容器下的 process 截圖 x1：

```
[dongrui@dongrui:~/os_hw3$ ps aux | grep 4457
root      4457   0.0  0.0  1632    4 pts/0    Ss+  10:39   0:00 sh
dongrui   4571   0.0  0.0  6476  2148 pts/4    S+   10:46   0:00 grep --color=
auto 4457
```

## 9. 心得:

(請寫下完成本次作業的心得、學到哪些東西、困難點的部分。)

我覺得這次的作業雖然有很多道步驟需要操作，但是基本上內容主要都是對比在 VM terminal、Docker container、src 資料夾裡 .c 檔的差異，或是觀察限制 CPU 後的使用率為何。從這次作業中，我可以發現到原來使用 namespace 能做到隔離的動

作，讓不同的終端機介面可以產生不同的 output，其中每個不一樣的 namespace（uts、ipc、pid、mnt）都有著各自的功用，對於開發者來說可以依照自身的需求去做搭配使用，感覺會是相當方便的使用手法。接著，對於 Docker 上使用資源限制與 Cgroup 的使用方法，我覺得可以限制 CPU 的使用率這一點算是很好的功能，因為有時候可能會希望某項東西不要佔用 CPU 太高的使用率，像這種時候就能使用從這次作業中學到的方法來限制 CPU 的使用率，非常好用。