

政大資科系

作業系統

Operating System

廖峻鋒

cfliao@nccu.edu.tw

Operating System

Deadlock

Chun-Feng Liao

廖峻鋒

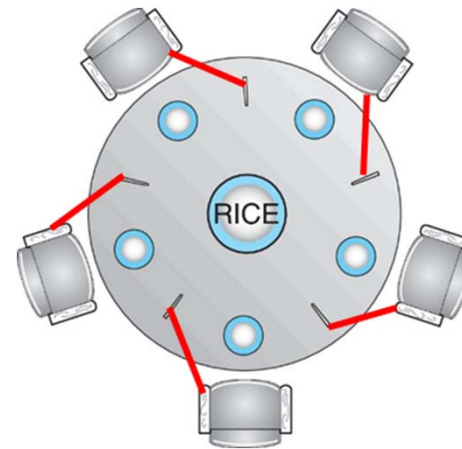
Department of Computer Science

National Chengchi University

Definition

- A set of threads is in a deadlocked state when
 - All thread in the set is waiting for an **event** that can be caused **only** by another thread in the set
 - Key events: resource acquisition and release

```
while (true) {  
    wait(chopstick[i]);  
    wait(chopstick[(i+1) % 5]);  
    . . .  
    /* eat for a while */  
    . . .  
    signal(chopstick[i]);  
    signal(chopstick[(i+1) % 5]);  
    . . .  
    /* think for awhile */  
    . . .  
}
```



Necessary Conditions

- 四個條件都**同時**成立，才會有deadlock
 - 免除其中一個條件就可避免deadlock
- deadlock**四要件**
 - **Mutual exclusion**: only one at a time can use a resource
 - **Hold & Wait**: someone holding some resources and is waiting for another resource
 - **No preemption**: a resource can only be released voluntarily
 - **Circular wait**: $\exists\{T_0, T_1, \dots, T_n\} \Rightarrow T_0 \rightarrow T_1 \rightarrow \dots \rightarrow T_n \rightarrow T_0$

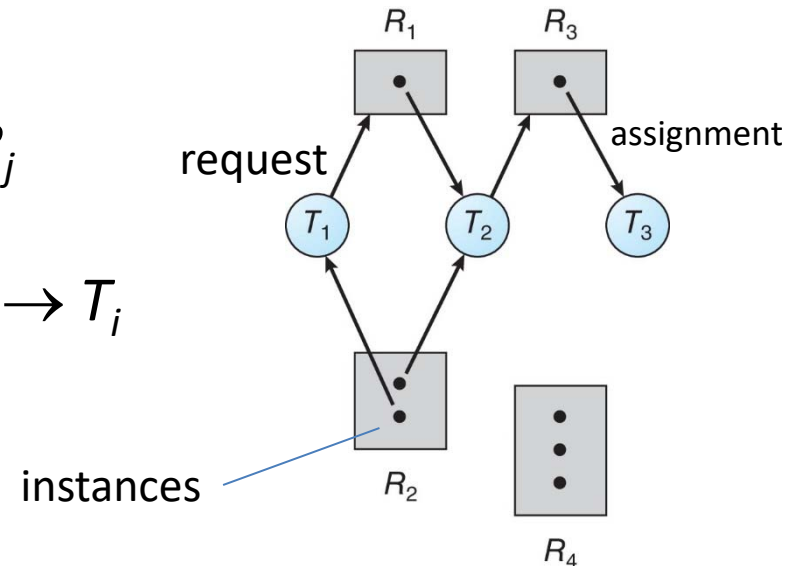
Resource-Allocation Graph

A set of vertices V and a set of edges E .

- V : two types
 - $T = \{T_1, T_2, \dots, T_n\}$, the set consisting of all the threads in the system
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system

- **request edge** – directed edge $T_i \rightarrow R_j$
- **assignment edge** – directed edge $R_j \rightarrow T_i$

If the graph contains a **cycle**, a deadlock **may exist!**



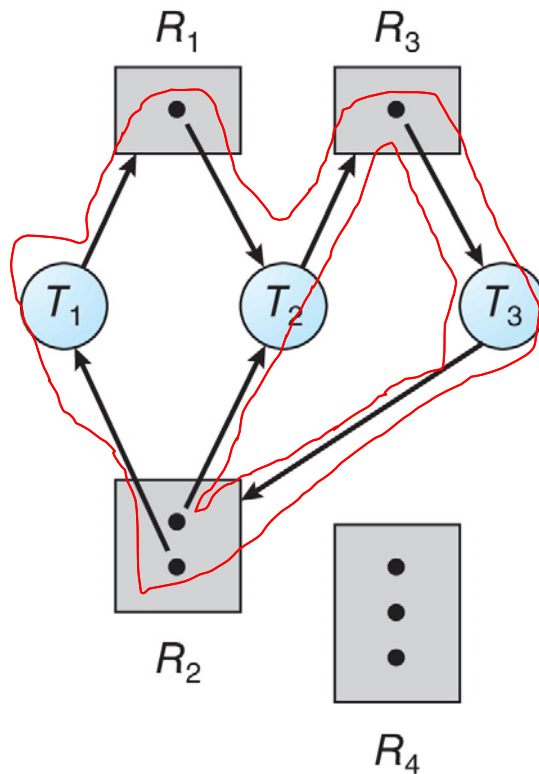
Basic Facts

- If graph contains no cycle → no deadlock
 - Circular wait cannot be held
- If graph contains a cycle:
 - if one instance per resource type → deadlock
 - if multiple instances per resource type → possibility of deadlock

Example

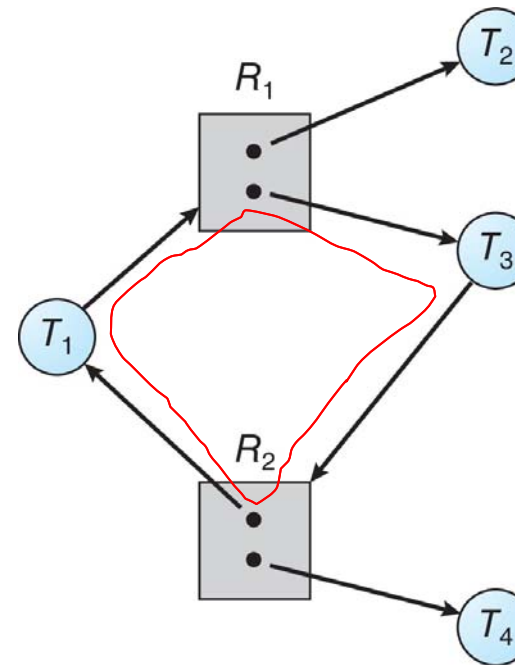
Deadlock: **All** threads in the set are waiting for **events** that can be caused **only** by another thread in the set

觀察重點: 先看有等資源的→評估是否可能等到



Deadlock

All threads are waiting for other threads



No deadlock

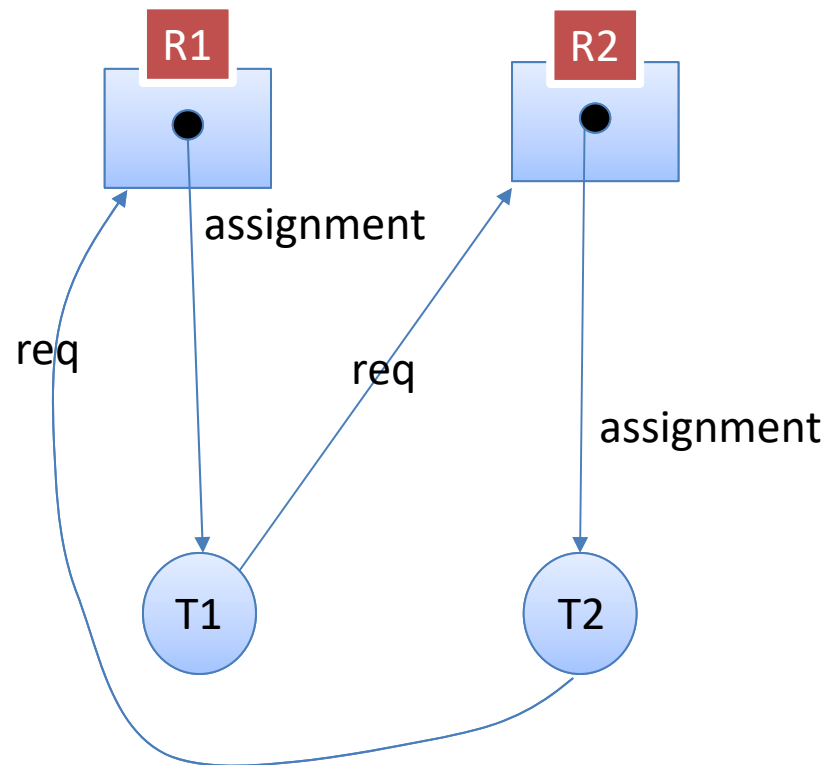
T_2, T_4 are not waiting for other threads

Example

T1

```
pthread_mutex_t first_mutex;
pthread_mutex_t second_mutex;

pthread_mutex_init(&first_mutex, NULL);
pthread_mutex_init(&second_mutex, NULL);
```



```
/* thread one runs in this function */
void *do_work_one(void *param)
{
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);

    pthread_exit(0);
}
```

T2

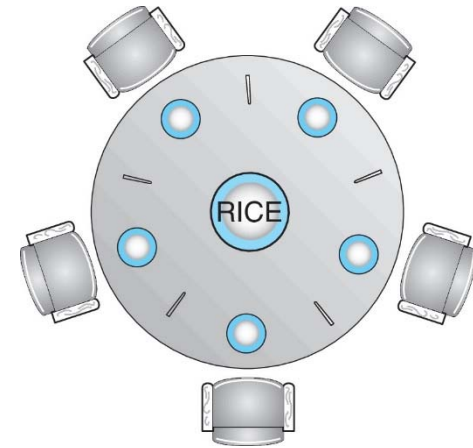
```
/* thread two runs in this function */
void *do_work_two(void *param)
{
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);

    pthread_exit(0);
}
```


Handling Deadlocks

- Ensure the system will never enter a deadlock state
 - **deadlock prevention**: 事前防範 ensure that at least one of the 4 necessary conditions cannot hold
 - **deadlock avoidance**: 邊走邊看 dynamically examines the resource-allocation state before allocation
- Allow deadlocks occur and then recover them
 - **deadlock detection**
 - **deadlock recovery**
- Ignore deadlocks
 - **used by most operating systems**, including LINUX and Windows

Deadlock Prevention



- Prevent “mutual exclusion”
 - Sharable resources
 - Ex: read-only files
 - Not always possible: some resources are not sharable
- Prevent “hold & wait”:
 - When requesting a resource, it does not hold any resource
 - Protocol1: Pre-allocate all resources before executing
 - Protocol2: Allow requesting resource when it has none
 - Resource utilization is low; starvation is possible

全放掉才
能拿

Deadlock Prevention

- Prevent “no preemption”
 - When a process is waiting on a resource, all its holding resources are preempted 在等資源的人所持有的資源隨時都可以被搶
 - e.g. T2 is waiting for R2 and holding R1 : R1 is preemptive
(T1 → R1 → T2 → R2) T2在等R2，所以它持有的R1會被preempted
R1 can be preempted and reallocated to T1
 - Applied to resources whose states can be easily saved and restored
 - e.g. CPU registers & memory
 - Cannot easily be applied to physical resources
 - e.g. printers & tape drives

Deadlock Prevention

- Prevent “Circular wait” Idea: 規定process只能按照一定的次序要求資源

- Impose a total ordering of all resources types
- A process requests resources in an increasing order

- Let $R = \{R_0, R_1, \dots, R_N\}$ be the set of resource types

$$R_i \rightarrow T_k \rightarrow R_j \Leftrightarrow F(R_i) < F(R_j) \quad F(x) \rightarrow x \text{ 的 order}$$

如果Tk已要求過Ri，現在要求Rj，則i一定小於j (亦即由小到大要求)

- Example:

- $F(\text{SSD drive}) = 1, F(\text{HDD drive}) = 5, F(\text{printer}) = 12$
- request order in program: SSD \rightarrow HDD \rightarrow printer

- Proof: assuming circular wait exists, then

先假設circular wait成立

$$T_0 \rightarrow R_0 \rightarrow T_1 \rightarrow R_1 \rightarrow T_2 \rightarrow R_2 \rightarrow \dots \rightarrow T_n \rightarrow R_n \rightarrow T_0 \rightarrow R_0$$

$$F(R_0) < F(R_1) < \dots < \underline{F(R_n) < F(R_0)} \Rightarrow \text{not possible!}$$

違反規定!

Order Example

T1

```
pthread_mutex_t first_mutex;  
pthread_mutex_t second_mutex;  
  
pthread_mutex_init(&first_mutex, NULL);  
pthread_mutex_init(&second_mutex, NULL);
```

```
/* thread one runs in this function */  
void *do_work_one(void *param)  
{  
    pthread_mutex_lock(&first_mutex);  
    pthread_mutex_lock(&second_mutex);  
    /**  
     * Do some work  
     */  
    pthread_mutex_unlock(&second_mutex);  
    pthread_mutex_unlock(&first_mutex);  
  
    pthread_exit(0);  
}
```

T2

```
/* thread two runs in this function */  
void *do_work_two(void *param)  
{  
    pthread_mutex_lock(&second_mutex);  
    pthread_mutex_lock(&first_mutex);  
    /**
```

T2要改成先要求first → then second
如此一來T1、T2只能有一個人先搶到first

```
    pthread_mutex_unlock(&second_mutex);  
  
    pthread_exit(0);  
}
```

Deadlock Avoidance

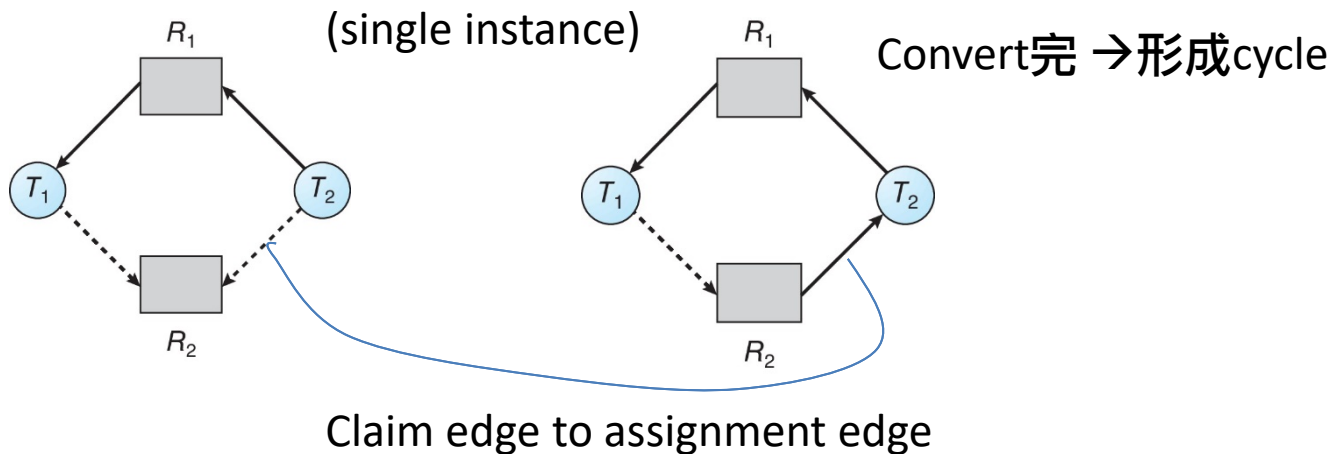
- 和Prevent的區別: 事前防 vs. 邊走邊閃
 - Prevent: Limit how request can be made (請求的方式)
 - Avoidance: Requires additional information about how resources are to be requested; 若認為有deadlock就reject請求
- Two approaches
 - Single instance resource
 - resource-allocation graph (RAG) algorithm
 - based on circle detection
 - Multiple instances resource
 - banker's algorithm based on safe sequence detection

RAG Algorithm

- Request edge: $T_i \rightarrow R_j$ T_i is waiting for R_j
- Assignment edge: $R_j \rightarrow T_i$ R_j has been allocated to T_i
- Claim edge: $T_i \dashrightarrow R_j$ T_i may request R_j in the future

- The request can be granted only if converting the claim edges to assignment edges does not result in cycles

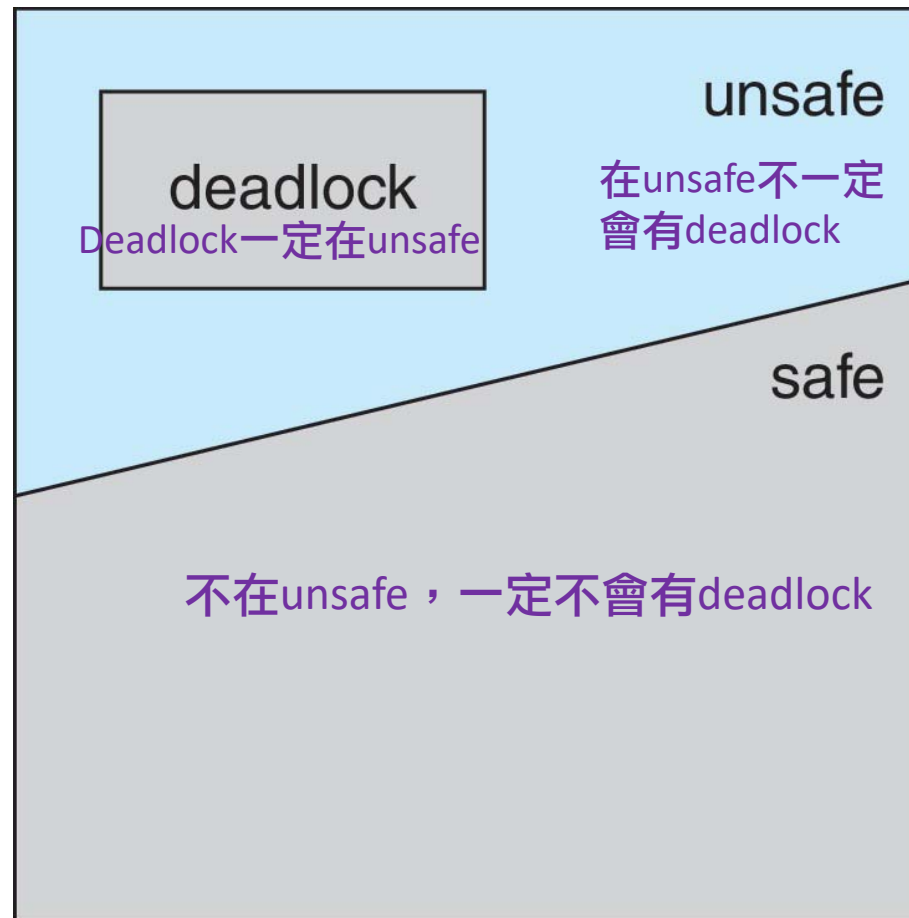
直覺意義: 如果未來R配置給T之後形成cycle的話，就不允許此要求



Limitations of RAG

- Single resource
- Resources must be claimed a priori in the system
 - 要事先規畫好，不能臨時多要
- Checking for safety using RAG is $O(n^2)$

Safe, Unsafe, Deadlock State



Safe State

- Safe state 存在至少一種滿足所有人需求的資源配置次序
 - A system is in a safe state if there exists a sequence of allocations to satisfy requests by all threads
 - This sequence of allocations is called safe sequence
- A system is in a safe state if there exists a sequence
Allocation sequence
 $\exists T = \langle T_1, T_2, \dots, T_n \rangle \rightarrow$
 $\forall T_i \in T, T_i$'s request can be satisfied by the free resources or
the resources held by T_j , where $j < i$
目前沒人用的 (現在可滿足)
j比i早: 有人用, 但接下來將釋出的 (未來可滿足)

Safe State with Safe Sequence

- There are 12 resources; 共12個，剩3個
- Assuming at t0: $\rightarrow \langle T1, T0, T2 \rangle$ is a safe sequence

一般都會考量 Worst case (Max needs)

Hint from
processes \rightarrow

Max Needs

Current Holding

T0

10

5

T1

4

2

(2) T1還需要2個

T2

9

2

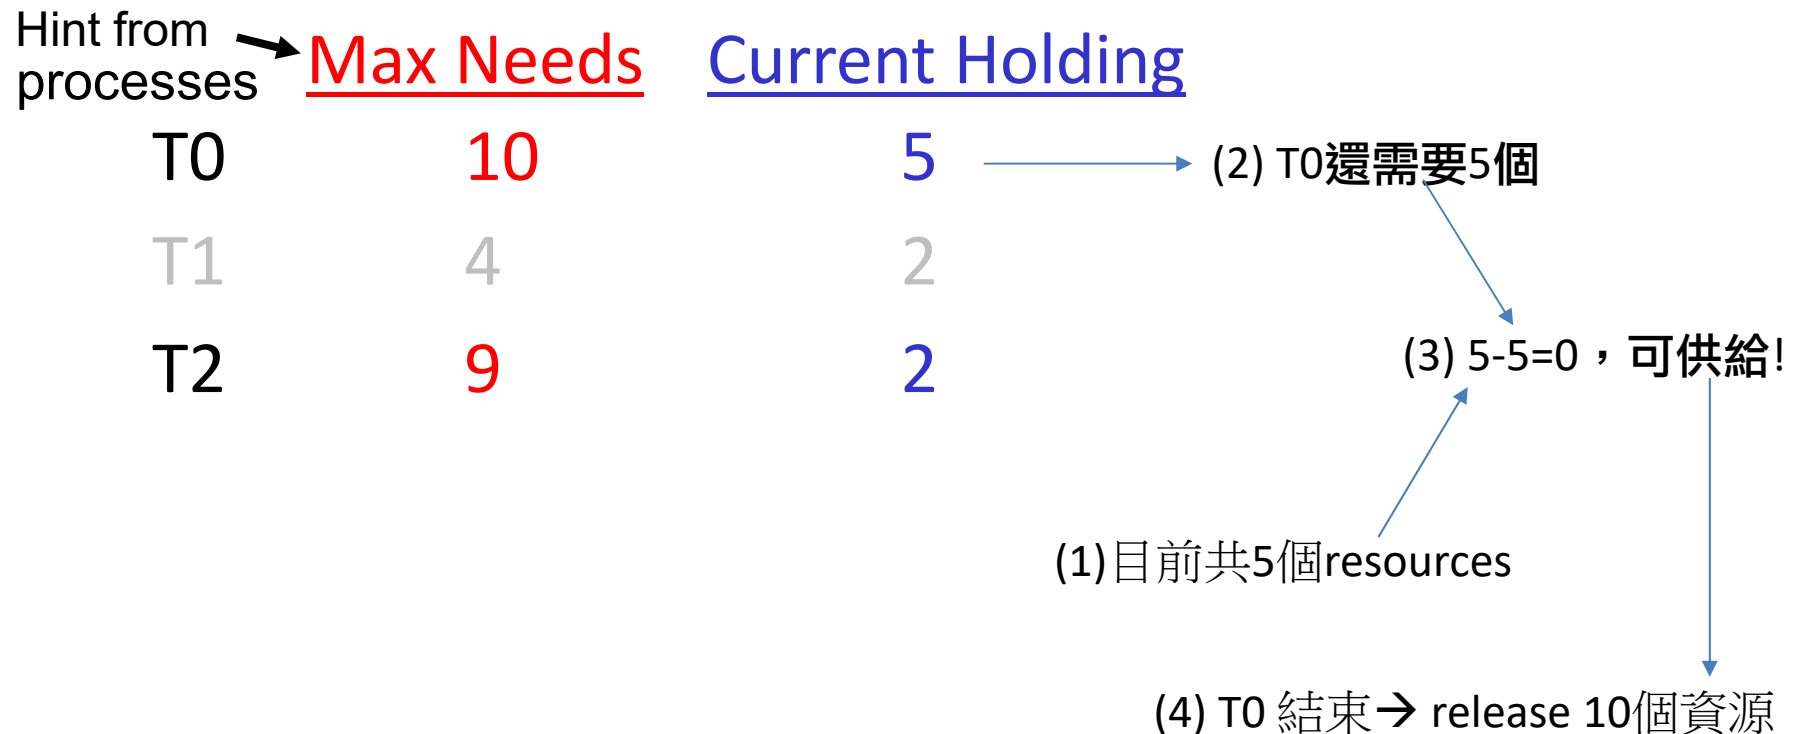
(3) $3-2=1$ ，可供給!

(1) 目前hold了9個resources; 剩 $12-9=3$ 個resources

(4) T1 結束 \rightarrow release 4個資源 $\rightarrow 1(\text{原來剩下的})+4 = 5$

Safe State with Safe Sequence

- 共12個，剩5個
- Assuming at t0: $\rightarrow \langle T1, T0, T2 \rangle$ is a safe sequence



Safe State with Safe Sequence

- 共12個，剩10個
- Assuming at t0: $\rightarrow \langle T1, T0, T2 \rangle$ is a safe sequence

Hint from
processes \rightarrow

Max Needs

Current Holding

T0 10

5

T1 4

2

T2 9

2

(2) 目前共10個resources

(3) $10 - 7 = 3$ ，可供給!

(1) T2還需要7個

(4) T2 結束 \rightarrow release 9個資源;
 $9 + 3 = 12$ 個資源

Un-Safe State without Safe Sequence

■ Assuming at t1:

	<u>Max Needs</u>	<u>Current Holding</u>	<u>Available</u>
T0	10	5	
T1	4	2	2
T2	9	2 3	

$12 - (5 + 2 + 3) = 2$

目前共hold了10個resources; $12 - 10 = 2$

T1 Max Needs 2 ($4 - 2 = 2$) $\rightarrow 2 - 2 = 0$ ，可供給

T1 結束 \rightarrow release 4 \rightarrow 總資源4

T0 Max Needs 5 > 4 ，無法供給!

T2 Max Needs 6 > 4 ，無法供給! (二條路都不通)

if T2 requests & is allocated 1 more resource

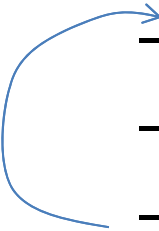
➔ 資源剩2，接下來只能選T1

➔ No safe sequence exist (T1 \rightarrow ... 二條路都不通)

➔ this allocation enters the system into an unsafe state

- A request is only granted if the allocation leaves the system in a safe state

Banker's Algorithm

- Use for multiple instances of each resource type
 - Banker algorithm:
 - A general safety algorithm to pre-determine if any safe sequence exists after allocation
 - Only proceed the allocation if safe sequence exists
 - Safety algorithm: P.335
 - Assume threads need **maximum** resources
 - Find a thread that can be satisfied by free resources
 - Free the resource usage of the thread
 - Goto step 2 (until all threads are satisfied)
- 

Banker's Algorithm Example (Safety Algo.)

- Total instances: A:10, B:5, C:7
- Used instances: A:7, B:2, C:5
- Available instances: A:3, B:3, C:2

	<u>Max</u>			<u>Allocation</u>		
	A	B	C	A	B	C
P0	7	5	3	0	1	0
P1	3	2	2	2	0	0
P2	9	0	2	3	0	2
P3	2	2	2	2	1	1
P4	4	3	3	0	0	2

(2) 7 2 5

Need(Max.-Alloc.)

(1)

A	B	C
7	4	3
1	2	2
6	0	0
0	1	1
4	3	1

3, 3, 2

綠色線代表可行

- Safe sequence: P1₍₄₎

Banker's Algorithm Example (Safety Algo.)

- Total instances: A:10, B:5, C:7
- Available instances: A:3, B:3, C:2

	<u>Max</u>			<u>Allocation</u>			<u>Need(Max.-Alloc.)</u>		
	A	B	C	A	B	C	A	B	C
P0	7	5	3	0	1	0	7	4	3
P1	3	2	2	2	0	0	1	2	2
P2	9	0	2	3	0	2	6	0	0
P3	2	2	2	2	1	1	0	1	1
P4	4	3	3	0	0	2	4	3	1

分配給P1資源並執行

3, 3, 2 原來剩下

1, 2, 2 P1還需

2, 1, 0 系統剩餘資源

完成後歸還資源

3, 2, 2 P1歸還

2, 1, 0 系統剩餘資源

5, 3, 2

- Safe sequence: P1

Banker's Algorithm Example (Safety Algo.)

- Total instances: A:10, B:5, C:7
- Available instances: A:5, B:3, C:2

	<u>Max</u>			<u>Allocation</u>			<u>Need(Max.-Alloc.)</u>		
	A	B	C	A	B	C	A	B	C
P0	7	5	3	0	1	0	7	4	3
P2	9	0	2	3	0	2	6	0	0
P3	2	2	2	2	1	1	0	1	1
P4	4	3	3	0	0	2	4	3	1

分配資源並執行

5, 3, 2 原來剩下

0, 1, 1 P3還需

5, 2, 1 系統剩餘資源

完成後歸還資源

2, 2, 2 P3歸還

5, 2, 1 系統剩餘資源

7, 4, 3

- Safe sequence: P1, P3

Banker's Algorithm Example (Safety Algo.)

- Total instances: A:10, B:5, C:7
- Available instances: A:7, B:4, C:3

	<u>Max</u>			<u>Allocation</u>			<u>Need(Max.-Alloc.)</u>		
	A	B	C	A	B	C	A	B	C
P0	7	5	3	0	1	0	7	4	3
P2	9	0	2	3	0	2	6	0	0
P4	4	3	3	0	0	2	4	3	1

- Safe sequence: P1, P3, P4

分配資源並執行

7, 4, 3 原來剩下

4, 3, 1 P4還需

3, 1, 2 系統剩餘資源

完成後歸還資源

4, 3, 3 P4歸還

3, 1, 2 系統剩餘資源

7, 4, 5

Banker's Algorithm Example (Safety Algo.)

- Total instances: A:10, B:5, C:7
- Available instances: A:7, B:4, C:5

	<u>Max</u>			<u>Allocation</u>			<u>Need(Max.-Alloc.)</u>		
	A	B	C	A	B	C	A	B	C
P0	7	5	3	0	1	0	7	4	3
P2	9	0	2	3	0	2	6	0	0

- Safe sequence: P1, P3, P4, P2

分配資源並執行

7, 4, 5 原來剩下

6, 0, 0 P2還需

1, 4, 5 系統剩餘資源

完成後歸還資源

9, 0, 2 P2歸還

1, 4, 5 系統剩餘資源

10, 4, 7

Banker's Algorithm Example (Safety Algo.)

- Total instances: A:10, B:5, C:7
- Available instances: A:10, B:4, C:7

	<u>Max</u>			<u>Allocation</u>			<u>Need(Max.-Alloc.)</u>		
	A	B	C	A	B	C	A	B	C
P0	7	5	3	0	1	0	7	4	3



- Safe sequence: P1, P3, P4, P2, P0

分配資源並執行

10, 4, 7 原來剩下

7, 4, 3 P0還需

3, 0, 4 系統剩餘資源

完成後歸還資源

3, 0, 4 P0歸還

7, 5, 3 系統剩餘資源

10, 5, 7

Resource Request

■ Total instances: A:10, B:5, C:7

■ Available instances: ~~A:3, B:3, C:2~~ → 2, 3, 0

	<u>Max</u>			<u>Allocation</u>			<u>Need(Max-Alloc)</u>		
	A	B	C	A	B	C	A	B	C
P0	7	5	3	0	1	0	7	4	3
→ P1	3	2	2	2	0	0	3, 0, 2	1	2
P2	9	0	2	3	0	2	6	0	0
P3	2	2	2	2	1	1	0	1	1
P4	4	3	3	0	0	2	4	3	1

■ If Request (P1) = (1, 0, 2): P1 allocation → 3, 0, 2

➤ Enter another safe state (Safe sequence: P1, P3, P4, P0, P2)

Resource Request

■ Total instances: A:10, B:5, C:7

■ Available instances: ~~A:3, B:3, C:2~~ $\xrightarrow{\text{A:多要3;B多要3}}$ 0, 0, 2

	<u>Max</u>			<u>Allocation</u>			<u>Need(Max-Alloc)</u>				
	A	B	C	A	B	C	A	B	C		
P0	7	5	3	0	1	0	7	4	3		
P1	3	2	2	2	0	0	1	2	2		
P2	9	0	2	3	0	2	6	0	0		
P3	2	2	2	2	1	1	0	1	1		
→ P4	4	3	3	0	0	2	→ 3, 3, 2	4	3	1	→ 1, 0, 1 = (4,3,3)-(3,3,2)

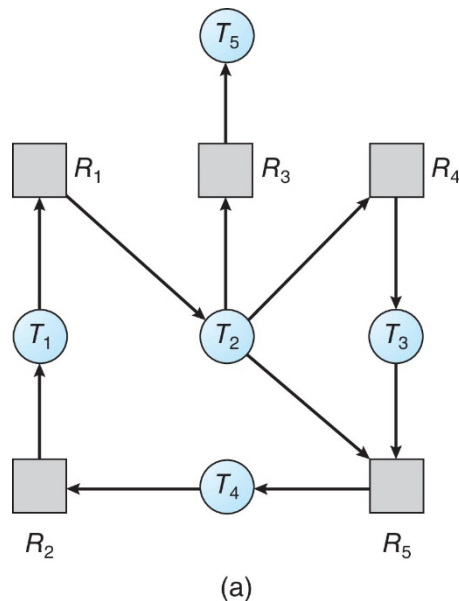
每個都不可行!

■ If Request (P4) = (3, 3, 0): P4 allocation \rightarrow 3, 3, 2

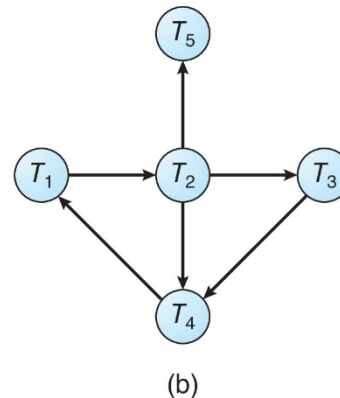
➤ enter into an unsafe state (no safe sequence can be found!)

Deadlock Detection

- Single instance 每個時刻都要做
 - convert request/assignment edges into wait-for graph
 - deadlock exists if there is a **cycle** in the wait-for graph



Resource-Allocation Graph



Corresponding wait-for graph

Multiple-Instance



P.340

- Total instances: A:7, B:2, C:6
- Available instances: A:0, B:0, C:0

	<u>Allocation</u>			<u>Request</u>		
	A	B	C	A	B	C
P0	0	1	0	0	0	0
P1	2	0	0	2	0	2
P2	3	0	3	0	0	0
P3	2	1	1	1	0	0
P4	0	0	2	0	0	2

- The system is in a safe state → $\langle P0, P2, P3, P1, P4 \rangle$
→ no deadlock
- If P2 request = $\langle 0, 0, 1 \rangle$ → no safe sequence can be found
→ the system is deadlocked

Deadlock Recovery

- Process termination
 - Abort all deadlocked processes
 - Abort 1 process at a time until the deadlock cycle is eliminated
 - which process should we abort first? (minimum cost)
- Resource preemption   Many factors, see P.342-343
 - Select a victim: which one to preempt?
 - Rollback: partial rollback or total rollback?
 - Starvation: can the same process be preempted always?

課後閱讀

- P.320 Livelock: 和Deadlock有什麼差別?
- P. 326 Database systems一般用什麼方法來解決 deadlock?
- P.342-343 many factors about process termination and resource preemption

Q & A